


Phase 3 sponsor summary: adjusted cost frontier (20260713)

Executive summary

- Scope: 15 valid full-suite arms, 900 benchmark attempts, and 515 raw successes.
- Best raw pass rate: **router-gpt-5.5** at **73.3%**, with adjusted known cost of **\$183.96**.
- Lowest cost per clean success among arms at or above 60% pass rate: **router-glm-5.1** at **\$0.58** per clean success.
- Recorded cost was **\$820.77**; adjusted known cost is **\$972.17**, revealing a known accounting gap of **\$151.40**.
- Failure/incomplete spend was **\$335.76** (34.5% of adjusted known cost). Unclean spend was **\$364.29** (37.5%).
- Failure/incomplete attempts consumed **213,482,015** input+output tokens (44.0%). Non-clean outcomes consumed **218,979,987** input+output tokens (45.1%).
- Remaining unresolved cost rows: **29**, all rows without usable cost or token metadata in the adjusted-cost layer.

Cost/performance frontier



The Pareto frontier by pass rate versus mean adjusted cost per benchmark attempt is: **router-glm-5.1**, **router-glm-5.2**, **router-anthropic-fable-5**, **router-deepeek-flash**, **router-gpt-5.5**. Because the full

Terminal-Bench suite uses three attempts per task, the chart uses cost per benchmark attempt. Multiply by three for a rough three-attempt-per-task budget.

Sponsor summary table

Arm	Pass rate	Adjusted cost	Cost / clean success	Failure/incomplete spend	Unclean spend	Unclean token share	Confidence
router-gpt-5.5	73.3%	\$183.96	\$4.38	\$41.64 (22.6%)	\$47.93 (26.1%)	24.9%	mixed
router-deepseek-flash	66.7%	\$56.80	\$1.42	\$13.88 (24.4%)	\$13.88 (24.4%)	61.8%	medium
router-anthropic-fable-5	65.0%	\$37.69	\$1.02	\$7.33 (19.4%)	\$17.38 (46.1%)	22.5%	medium
router-anthropic-opus	65.0%	\$64.49	\$1.65	\$30.74 (47.7%)	\$30.74 (47.7%)	21.8%	mixed
router-gpt-5.4	65.0%	\$183.65	\$4.83	\$34.49 (18.8%)	\$39.60 (21.6%)	20.1%	low
router-glm-5.2	63.3%	\$25.32	\$0.72	\$6.87 (27.1%)	\$8.23 (32.5%)	57.5%	medium
router-gemini-3.1-pro	63.3%	\$46.47	\$1.29	\$13.82 (29.7%)	\$14.46 (31.1%)	25.3%	low
router-deepseek-pro	61.7%	\$50.44	\$1.36	\$12.29 (24.4%)	\$12.29 (24.4%)	26.8%	medium
router-glm-5.1	60.0%	\$20.30	\$0.58	\$3.94 (19.4%)	\$4.10 (20.2%)	31.0%	mixed
router-grok-build-0.1	60.0%	\$53.15	\$1.48	\$29.83 (56.1%)	\$29.83 (56.1%)	54.9%	mixed
router-qwen-3.7-plus	55.0%	\$34.94	\$1.09	\$16.47 (47.1%)	\$19.16 (54.8%)	53.2%	mixed
router-kimi-k2.6	50.0%	\$35.13	\$1.17	\$22.07 (62.8%)	\$22.07 (62.8%)	56.2%	low

Arm	Pass rate	Adjusted cost	Cost / clean success	Failure/incomplete spend	Unclean spend	Unclean token share	Confidence
router-anthropic-sonnet	46.7%	\$52.79	\$1.89	\$23.32 (44.2%)	\$23.32 (44.2%)	40.4%	mixed
router-anthropic-haiku-sanitized	41.7%	\$83.51	\$3.34	\$54.82 (65.6%)	\$54.82 (65.6%)	65.7%	high
router-gemini-flash	21.7%	\$43.53	\$4.84	\$24.26 (55.7%)	\$26.47 (60.8%)	56.7%	mixed

Value-oriented shortlist

Arm	Pass rate	Adjusted cost	Cost / clean success	Failure/incomplete spend	Unclean spend	Unclean token share	Confidence
router-glm-5.1	60.0%	\$20.30	\$0.58	\$3.94 (19.4%)	\$4.10 (20.2%)	31.0%	mixed
router-glm-5.2	63.3%	\$25.32	\$0.72	\$6.87 (27.1%)	\$8.23 (32.5%)	57.5%	medium
router-anthropic-fable-5	65.0%	\$37.69	\$1.02	\$7.33 (19.4%)	\$17.38 (46.1%)	22.5%	medium
router-gemini-3.1-pro	63.3%	\$46.47	\$1.29	\$13.82 (29.7%)	\$14.46 (31.1%)	25.3%	low
router-deepseek-pro	61.7%	\$50.44	\$1.36	\$12.29 (24.4%)	\$12.29 (24.4%)	26.8%	medium
router-deepseek-flash	66.7%	\$56.80	\$1.42	\$13.88 (24.4%)	\$13.88 (24.4%)	61.8%	medium
router-grok-build-0.1	60.0%	\$53.15	\$1.48	\$29.83 (56.1%)	\$29.83 (56.1%)	54.9%	mixed

Arm	Pass rate	Adjusted cost	Cost / clean success	Failure/incomplete spend	Unclean spend	Unclean token share	Confidence
router-anthropic-opus	65.0%	\$64.49	\$1.65	\$30.74 (47.7%)	\$30.74 (47.7%)	21.8%	mixed

Cost accounting and nonproductive spend

- The original recorded-cost view understated known spend by **\$151.40**, or **18.4%** over recorded cost.
- The gap is concentrated in errored or operationally unclean paths, which is why adjusted cost is more appropriate for sponsor-facing cost comparisons.
- Failure/incomplete spend directly measures money spent on attempts that did not pass: **\$335.76**.
- Unclean spend includes failures, incomplete outcomes, and exception-with-success-signal rows: **\$364.29**.
- Token-volume waste is directionally similar: failure/incomplete attempts consumed **213,482,015** tokens; all non-clean outcomes consumed **218,979,987** tokens.

Method note

- Recorded cost remains preserved as imported artifact truth.
- Adjusted known cost adds reconstructed missing-cost rows using configured pricing snapshots or same-arm empirical estimates.
- Invalid/quarantined runs are excluded from this valid-only comparison.
- Token-volume analysis uses input + output tokens. Cache tokens affect cost but are not added to the token-volume numerator to avoid double-counting cached input.
- Provider invoices/dashboards remain separate evidence; this report is a benchmark-level adjusted-cost view.

Generated artifacts

- Frontier chart:
[../../../../results/phase3/reporting/phase3_adjusted_cost_frontier_20260713.svg](#)
- Sponsor summary table:
[results/phase3/reporting/phase3_sponsor_summary_table_20260713.tsv](#)
- Token/outcome breakdown:
[results/phase3/reporting/phase3_token_outcome_breakdown_20260713.tsv](#)